

# Lecture 2 Deep Feedforward Networks

**课程：机器学习与深度学习**

# Three Steps for Deep Learning



Deep Learning is as simple as linear model.....


# Overview

- Model Architectures
  - Artificial neurons
  - Activation function and saturation
  - Feedforward neural nets
- How to train a neural net
  - Loss Function Design
  - Optimization
    - Gradient Descent and Stochastic Gradient Descent
    - Back-propagation
  - Advanced Training tips

# The Perceptron

- Invented in 1954 by Frank Rosenblatt
- Inspired by neurobiology

# Artificial Neuron




- Each neuron is a very simple function
- Pre-activation:  $z(x) = \sum_i w_i x_i + b = w^T x + b$
- Output activation:  $a(x) = f(z(x)) = f(w^T x + b)$

$f(\cdot)$ : nonlinear activation function

$w$ : weight

$b$ : bias term




# Artificial Neuron

- Output activation

$$a(x) = f(z(x)) = f(w^T x + b)$$

Range is  
determined  
by  $g(\cdot)$

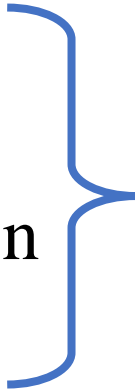


(from Pascal Vincent's slides)

Bias only changes  
the position of the  
riff

# Activation Function

- Linear activation function
- Sigmoid activation function
- Hyperbolic tangent activation function
- Rectified linear (ReLU) activation function
- Softmax activation function



Non-linear  
activation  
function,  
frequently  
used in  
deep neural  
networks.

# Linear Activation Function

- No input squashing
- No nonlinear transformation

$$f(z) = z$$


- Why non-linearity?
  - **Without non-linearity**, deep neural networks work the same as linear transform
    - $W_1(W_2 \cdot x) = (W_1 W_2)x = Wx$
  - **With non-linearity**, networks with more layers can approximate more complex function



# Sigmoid Activation Function

- Squashes the neuron's output to (0,1)
- Bounded
- Strictly increasing


$$f(z) = \sigma(z) \stackrel{\text{def}}{=} \frac{1}{1 + \exp(-z)}$$



# Hyperbolic Tangent (“tanh”) Function

- Squashes the neuron’s output to (-1,1)
- Can be positive or negative
- Bounded
- Strictly increasing


$$f(z) = \tanh(z) \stackrel{\text{def}}{=} \frac{\exp(2z) - 1}{\exp(2z) + 1}$$



# Rectified Linear Activation Function (ReLU)


- Tends to produce units with sparse activities
- No upper-bound
- Increasing

$$f(z) = \text{reclin}(z) \stackrel{\text{def}}{=} \max(z, 0)$$



- ReLU variants:

- Shift ReLU:  $\max(-1, z)$
- Leaky ReLU:  $\max(0.1z, z)$
- Parameter ReLU:  $\max(\mu z, z)$
- Exponential Linear Units:  $\max(z, \mu(\exp(z) - 1))$
- Maxout:  $\max(w_1^T z + b_1, w_2^T z + b_2)$



# Softmax Activation Function

- In multi-class classification ( $C$  classes), we need to
  - generate multiple output:  $\mathbf{z} \in \mathbb{R}^C$  (in vector form)
  - estimate the conditional probability of each class:



$$p(y = i|\mathbf{z}) = \frac{\exp z_i}{\sum_c \exp z_c}$$

- Strictly positive
- Sum up to one

$$\mathbf{f}(\mathbf{z}) = \text{softmax}(\mathbf{z}) = \left[ \frac{\exp z_1}{\sum_c \exp z_c} \cdots \frac{\exp z_C}{\sum_c \exp z_c} \right]^T$$


# Capacities of a Single Neuron

- Solve linearly separable problems




# Capacities of a Single Neuron


- Can't solve linearly inseparable problems



# How to implement XOR?




Input		Output
A	B	
0	0	0
0	1	1
1	0	1
1	1	0



$$A \text{ XOR } B = AB' + A'B$$

Multiple operations can produce more complicate output

# Neural Network



## Neural Network


Different connection leads to different network structures

Network parameter  $\theta$ : all the weights and biases in the “neurons”




# Multilayer Feedforward Neural Networks

- Each neuron in one layer has directed connections to the neurons of the subsequent layer
- Information propagates from input  $x$  to output  $f(x)$ , through many hidden layers




# Expressions of Multi-Layer Neural Network

Continuous function w/ 2 layers




Continuous function w/ 3 layers




Multiple layers enhance the model expression  $\rightarrow$  the model can approximate more complex functions

# Setting the Number of Neurons and Layers

- More neurons = more capacity
- More layers = more capacity




# Fully Connect Feedforward Network



# Example Application



## Input




16 x 16 =  
256

Ink  $\rightarrow$  1

No ink  $\rightarrow$  0

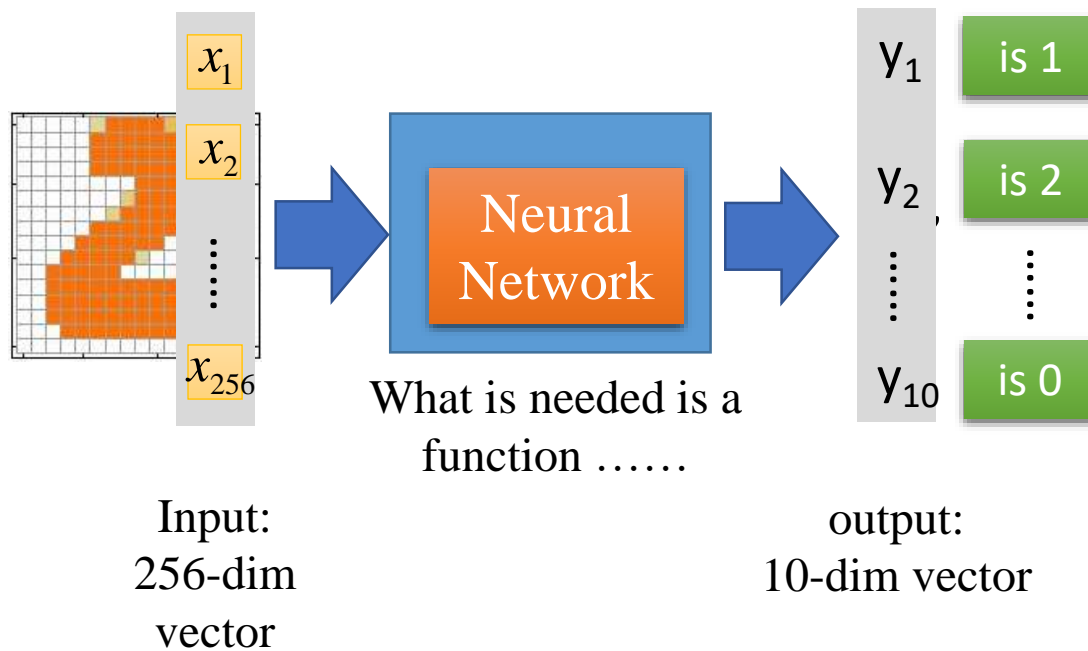
## Output



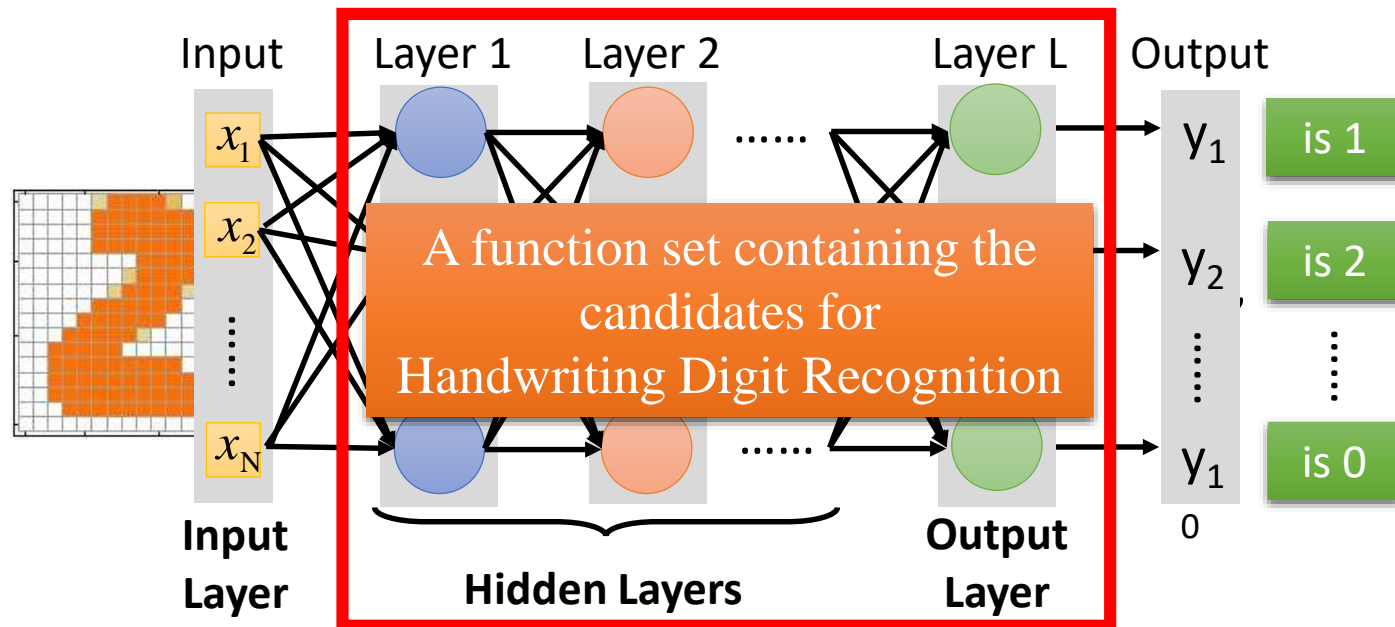
Each dimension  
represents the confidence  
of a digit.

# Example Application

- Handwriting Digit Recognition




# Example Application



You need to decide the network structure to let a good function in your function set.

# FAQ



- Q: How many layers? How many neurons for each layer?

Trial and Error

+



Intuition

- Q: Can the structure be automatically determined?
  - E.g. Evolutionary Artificial Neural Networks
- Q: Can we design the network structure?

Convolutional Neural Network (CNN)



# Notation Definition



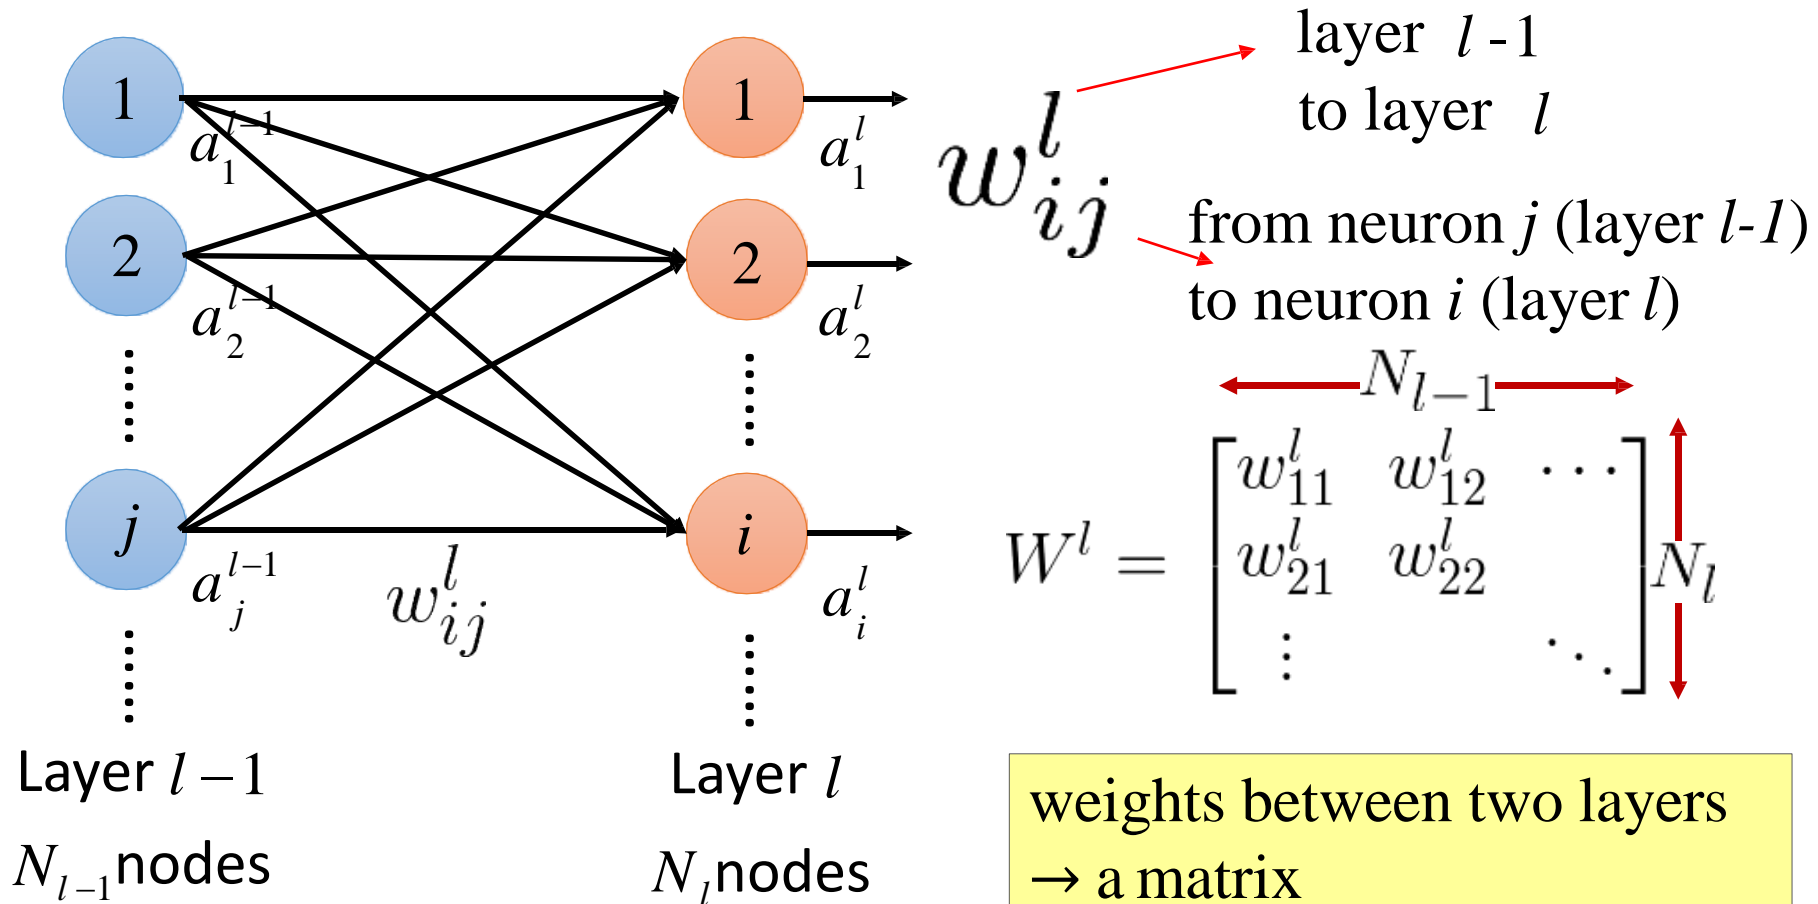
Output of a neuron:

$a_i^l$   $\xrightarrow{\text{red arrow}}$  layer  $l$   
 $a_i^l$   $\xrightarrow{\text{red arrow}}$  neuron  $i$


$$a^l = \begin{bmatrix} \vdots \\ a_i^l \\ \vdots \end{bmatrix}$$

output of one layer  $\rightarrow$  a vector

# Notation Definition



# Notation Definition

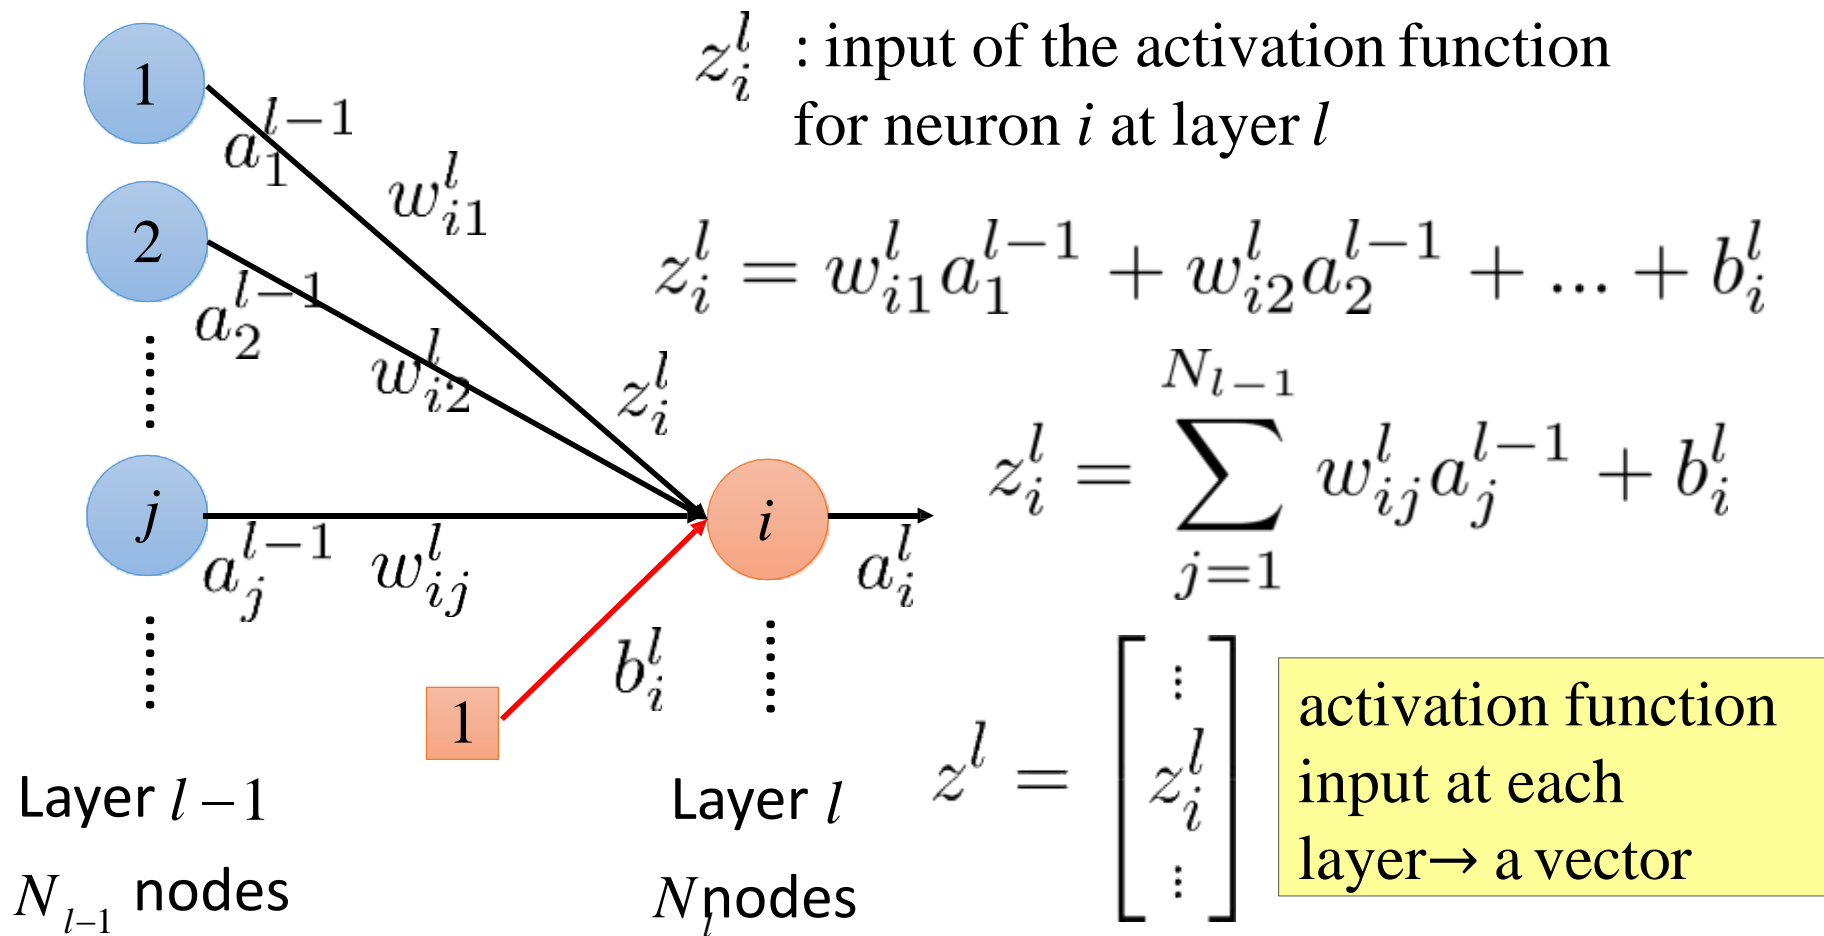


$b_i^l$  : bias for neuron  $i$   
at layer  $l$

$$b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

bias of all neurons at each  
layer  $\rightarrow$  a vector

# Notation Definition



# Notation Summary

---

$a_i^l$  : output of a neuron

$w_{ij}^l$  : a weight

$a^l$  : output vector of a layer

$W^l$  : a weight matrix

$z_i^l$  : input of activation  
function

$b_i^l$  : a bias


$z^l$  : input vector of activation  
function for a layer

$b^l$  : a bias vector

# Neural Network Formulation

$$f : R^N \rightarrow R^M$$

Fully connected feedforward network



$$\sigma(W^1 x + b^1) \parallel \sigma(W^2 a^1 + b^2) \parallel \sigma(W^L a^{L-1} + b^L) = y$$

$$y = f(x) = \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L$$

# Three Steps for Deep Learning



# Function = model parameters

**Forward propagation**

$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

Different parameters  $W$  and  $b \rightarrow$  different functions

Formal definition:

$$f(x; \theta); \theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

Pick a function  $f$  = pick a set of model parameters  $\theta$



# Training

- Empirical risk minimization  $J(\theta)$

$$\arg \min_{\theta} \frac{1}{T} \sum_t \underbrace{l(f(\mathbf{x}^{(t)}; \theta), y^{(t)})}_{\text{Loss function}} + \underbrace{\lambda \Omega(\theta)}_{\text{Regularizer}}$$

- Learning is cast as optimization
  - Find a model parameter set that minimize  $J(\theta)$
  - Loss function can sometimes be viewed as **a surrogate for what we want to optimize**

# Loss Function

- In discriminative model (判别模型), model  $y|x$ .
- Learning the maximum likelihood, equivalently the cross entropy between training data and model distribution:

$$l(\theta) = -E_{x,y \sim \text{Data}} \log p(y|x)$$

- The specific form of loss function depends on the model distribution  $p(\cdot)$

# Loss Functions

- Loss function evaluates the performance of our model, it is chosen according to the output units
  - Normal:  $\hat{y} = \mathbf{w}^T \mathbf{x} + b$
  - Bernoulli:  $\hat{y} = \sigma(\mathbf{w}^T \mathbf{x} + b)$
  - Multinomial:  $\hat{\mathbf{y}} = \text{softmax}(\mathbf{W}^T \mathbf{x} + \mathbf{b})$
- Consider regularization  $\Omega(\theta)$
- Equivalent form of Loss function:  $J = l(y, \hat{y}) + \lambda \Omega(\theta)$

# Frequently Used Loss Functions

- Square loss

$$l(y, \hat{y}) = (y - \hat{y})^2$$

- Hinge loss

$$l(y, \hat{y}) = \max(0, 1 - \hat{y} y)$$

- Logistic loss

$$l(y, \hat{y}) = \log(1 + \exp(-\hat{y} y))$$

- Cross entropy loss

$$l(y, \hat{y}) = -y \log \hat{y} + (1 - y) \log(1 - \hat{y})$$

# How to Train Multilayer Neural Nets?

- Learning is reduced to optimization.
  - Given a loss function  $J(\theta)$  and several parameter sets
  - Find a model parameter set that minimize  $J(\theta)$


# Overview

- Model Architectures
  - Artificial neurons
  - Activation function and saturation
  - Feedforward neural nets
- How to train a neural net
  - Loss Function Design
  - Optimization
    - Gradient Descent and Stochastic Gradient Descent
    - Backward propagation

# Gradient Descent for Optimization


Assume that  $\theta$  has only one variable

$\theta^i$  : the model at the i-th iteration




# Gradient Descent for Optimization

Assume that  $\theta$  has only one variable





# Gradient Descent for Optimization



## Algorithm

Initialization: start at  $\theta^0$   
while( $\theta^{(i+1)} \neq \theta^i$ )


{

compute gradient at  $\theta^i$   
update parameters

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$

}

# Revisit Neural Network Formulation



$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

# Gradient Descent for Neural Network

$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \cdots \\ w_{21}^l & w_{22}^l & \\ \vdots & & \ddots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla J(\theta) = \begin{bmatrix} \vdots \\ \frac{\partial J(\theta)}{\partial w_{ij}^l} \\ \vdots \\ \frac{\partial J(\theta)}{\partial b_i^l} \\ \vdots \end{bmatrix}$$

## Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{

    compute gradient at  $\theta^i$

    update parameters

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$


}

# Gradient Descent for Optimization

## Simple Case

$$y = f(x; \theta) = \sigma(Wx + b)$$

$$\theta = \{W, b\} = \{w_1, w_2, b\}$$



$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \frac{\partial J(\theta)}{\partial w_1} \\ \frac{\partial J(\theta)}{\partial w_2} \\ \frac{\partial J(\theta)}{\partial b} \end{bmatrix}$$

### Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{

compute gradient at  $\theta^i$


update parameters

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$

}

$$\begin{bmatrix} w_1^{i+1} \\ w_2^{i+1} \\ b^{i+1} \end{bmatrix} = \begin{bmatrix} w_1^i \\ w_2^i \\ b^i \end{bmatrix} - \eta \begin{bmatrix} \frac{\partial J(\theta^i)}{\partial w_1} \\ \frac{\partial J(\theta^i)}{\partial w_2} \\ \frac{\partial J(\theta^i)}{\partial b} \end{bmatrix}$$

# To compute the Gradients



- If square loss

- $\hat{y} = \sigma(Wx + b) = \sigma(w_1x_1 + w_2x_2 + b)$

- $J(\theta) = (\sigma(Wx + b) - y)^2$

- $\frac{\partial J(\theta)}{\partial w_1} = 2(\sigma(Wx + b) - y)(1 - \sigma(Wx + b))\sigma(Wx + b)x_1$

- $\frac{\partial J(\theta)}{\partial w_2} = 2(\sigma(Wx + b) - y)(1 - \sigma(Wx + b))\sigma(Wx + b)x_2$

- $\frac{\partial J(\theta)}{\partial b} = 2(\sigma(Wx + b) - y)(1 - \sigma(Wx + b))\sigma(Wx + b)$

## Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{compute gradient at  $\theta^i$

update parameters

$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$  }

# Gradient Descent for Neural Network

$$y = f(x) = \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \dots \\ w_{21}^l & w_{22}^l & \dots \\ \vdots & & \ddots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla J(\theta) = \begin{bmatrix} \vdots \\ \frac{\partial J(\theta)}{\partial w_{ij}^l} \\ \vdots \\ \frac{\partial J(\theta)}{\partial b_i^l} \\ \vdots \end{bmatrix}$$

## Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{

    compute gradient at  $\theta^i$


    update parameters

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$

}


# Gradient Computation: Sigmoid Unit

$$f(z) = \sigma(z), \quad f'(z) = \sigma(z)(1 - \sigma(z))$$



# Gradient Computation: Tanh Unit

- $f(z) = \tanh(z) = 2\sigma(2z) - 1$  and  $\tanh'(z) = 1 - f(z)^2$
- $\tanh(z)$  approximates linear function when  $z$  is small
- Often it is preferable to sigmoid in feedforward neural nets (zero-centered)
- Problem: still kill gradient when saturated ☹




[LeCun et al., 1991]



# Gradient Computation: ReLU

- $f(z) = \max(z, 0)$   
 $f'(z) = \begin{cases} 1, & z \geq 0 \\ 0, & \text{o. w.} \end{cases}$
- Models are easier to optimize if their behavior is close to linear
- Converge much faster than sigmoid and tanh in practice (6x faster)
- Not differentiable at  $z = 0$ , but it is not a problem in practice
- Not zero-centered
- Fragile during training and can “die”



[Krizhevsky et al., 2012]

# Gradient Computation: Leaky ReLU

- $f(a) = \max(z, 0.01z)$  `
- $f'(a) = \begin{cases} 1, & z \geq 0 \\ 0.01, & \text{o. w.} \end{cases}$


- Will not die

- Parametric ReLU

- $f(z) = \max(z, \mu z)$

- $f'(z) = \begin{cases} 1, & z \geq 0 \\ \mu, & \text{o. w.} \end{cases}$

- Update  $\mu$  through backpropagation




[Mass et al., 2013]

[He et al., 2015]

# Gradient Computation: Exponential Linear Unit (ELU)

- $f(z) = \begin{cases} z, & z \geq 0 \\ \mu(e^z - 1) & \end{cases}$
- $f'(z) = \begin{cases} 1, & z \geq 0 \\ \mu e^z, & \text{o.w.} \end{cases}$
- All benefit of ReLU
- Almost zero-centered
- Compute  $\exp()$  ☹️



[Clevert et al., 2015]

# Gradient Computation: Maxout

- $f(z) = \max(w_1 z + b_1, w_2 z + b_2)$
- Generalization of Leaky ReLU and ReLU
- Double the number of parameters

# Softmax

- Cross-entropy:

$$H(p, q) = - \sum_x p(x) \log(q(x))$$

- In multiclass classification, <sup>y-th</sup>

$$\mathbf{y} = [0, 0, 0, \dots, 0, 1, 0, 0, \dots 0]^T \in \mathbb{R}^C,$$

Then

$$J = H(\mathbf{y}, \mathbf{h}) = -\log h_y$$

where  $h_i = f_i(\mathbf{z}) = \frac{\exp(z_i)}{\sum_c^C \exp(z_c)} = \text{P}(y = i | \mathbf{z}), 1 \leq i \leq C.$

# Softmax

- $\nabla_{\mathbf{h}} J = [0, \dots, 0, -h_y, 0, \dots, 0]^T$
- $\nabla_{\mathbf{z}} J = \left( \frac{\partial \mathbf{h}}{\partial \mathbf{z}} \right)^T \nabla_{\mathbf{h}} J = -\frac{1}{h_y} \nabla h_y(\mathbf{z}) = -(\mathbf{e}_y - \mathbf{h})$

# In Practice

- For forward nn
  - Use ReLU, be careful with the learning rate
  - Tryout Leaky ReLU, ELU and Maxout
  - Tryout Tanh with low expectation
  - Never use sigmoid

# Gradient Descent for Neural Network

$$y = f(x) = \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \dots \\ w_{21}^l & w_{22}^l & \dots \\ \vdots & & \ddots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla J(\theta) = \left[ \dots, \frac{\partial J(\theta)}{\partial w_{ij}^l}, \dots, \frac{\partial J(\theta)}{\partial b_i^l}, \dots \right]^T$$

## Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{

    compute gradient at  $\theta^i$

    update parameters

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$

}

To compute the gradients of millions of parameters efficiently, we use **backpropagation**.



# Gradient Descent Issue

- After see all training samples, the model can be updated **slowly**.
- It is too expensive to compute the full gradient

Thus, we have stochastic gradient descent (SGD)

# Stochastic Gradient Descent

For  $t = 1, 2, 3, \dots$  (**epoch** means one pass over the full training set)

sample  $i \in \{1, 2, \dots, n\}$

$$\mathbf{s} = \nabla l(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t)}), y^{(i)}) + \lambda \nabla \Omega(\boldsymbol{\theta}^{(t)})$$


$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta \cdot \mathbf{s}$$

- Computing stochastic gradient is much cheaper than full gradient

# Gradient Descent v.s. SGD


## Gradient Descent

Update after seeing all examples



## Stochastic Gradient Descent

If there are 20 examples, update 20 times in one epoch.



SGD approaches to the target point faster than gradient descent

# Mini-Batch SGD

---

**Algorithm 8.1** Stochastic gradient descent (SGD) update at training iteration  $k$

---

**Require:** Learning rate  $\epsilon_k$ .

**Require:** Initial parameter  $\theta$

**while** stopping criterion not met **do**

    Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

    Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow +\frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$


    Apply update:  $\theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$

**end while**

---



# Handwriting Digit Classification

Training speed: mini-batch > SGD > Gradient Descent



# Big Issue: Local Optima

- Neural networks has no guarantee for obtaining global optimal solution
- Saddle points



**Advanced practical tips (to be presented in the last lecture)**

# Summary: How to Train Multilayer Neural Nets?

- Define the loss function  $l(,)$  properly
- A procedure to compute loss  $l(,)$  (*forward propagation*)
- A procedure to compute gradient  $\nabla l(,)$  (*back propagation*)
- Regularizer and its gradient  $\Omega(,)$  and  $\nabla \Omega(,)$
- Perform gradient based optimization method

# Forward/Backward Propagation

```
class ComputationalGraph(object):  
    #...  
    def forward(inputs):  
        # 1. [pass inputs to input gates...]  
        # 2. forward the computational graph:  
        for gate in self.graph.nodes_topologically_sorted():  
            gate.forward()  
        return loss # the final gate in the graph outputs the loss  
    def backward():  
        for gate in reversed(self.graph.nodes_topologically_sorted()):  
            gate.backward() # little piece of backprop (chain rule applied)  
        return inputs_gradients
```




# Overview

- Model Architectures
  - Artificial neurons
  - Activation function and saturation
  - Feedforward neural nets
- How to train a neural net
  - Loss Function Design
  - Optimization
    - Gradient Descent and Stochastic Gradient Descent
    - Back-propagation

# Forward v.s. Back Propagation

- In a feedforward neural network
  - forward propagation
    - from input  $x$  to output  $y$  information flows forward through the network
    - during training, forward propagation can continue onward until it produces a scalar cost  $J(\theta)$
  - back-propagation
    - allows the information from the cost to then flow backwards through the network, in order to compute the **gradient**
    - can be applied to any function



# Chain Rule

$$\Delta w \rightarrow \Delta x \rightarrow \Delta y \rightarrow \Delta z$$


$$\frac{\partial z}{\partial w} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w}$$

$$= f'(y) f'(x) f'(w)$$

forward propagation for loss (cost)

$$= f'(f(f(w))) f'(f(w)) f'(w)$$

back-propagation for gradient



# Gradient Descent for Optimization

$$y = f(x) = \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \dots \\ w_{21}^l & w_{22}^l & \dots \\ \vdots & & \ddots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla J(\theta) = \begin{bmatrix} \vdots \\ \frac{\partial J(\theta)}{\partial w_{ij}^l} \\ \vdots \\ \frac{\partial J(\theta)}{\partial b_i^l} \\ \vdots \end{bmatrix}$$

## Algorithm

Initialization: start at  $\theta^0$

while( $\theta^{(i+1)} \neq \theta^i$ )

{


    compute gradient at  $\theta^i$

    update parameters

}


$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} J(\theta^i)$$

$$\frac{\partial J(\theta)}{\partial w_{ij}^l}$$



$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} \quad (l > 1)$$




$$z^l = W^l a^{l-1} + b^l$$


$$z_i^l = \sum_j w_{ij}^l a_j^{l-1} + b_i^l$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} = a_j^{l-1}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} \quad (l = 1)$$



$$\frac{\partial J(\theta)}{\partial w_{ij}^l}$$




$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} = \begin{cases} a_j^{l-1} & , l > 1 \\ x_j & , l = 1 \end{cases}$$



$$\frac{\partial J(\theta)}{\partial w_{ij}^l}$$




$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial J(\theta)}{\partial z_i^l}$$

$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$\delta_i^l$  : the propagated gradient  
corresponding to the  $l$ -th layer



Idea: computing  $\delta^l$  layer by layer (from  $\delta^L$  to  $\delta^1$ ) is more efficient

$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

**Idea:** from L to 1

- (1) Initialization: compute  $\delta^L$
- (2) Compute  $\delta^l$  based on  $\delta^{l+1}$

$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

Idea: from L to 1

**(1) Initialization: compute  $\delta^L$**

(2) Compute  $\delta^l$  based on  $\delta^{l+1}$

$$\delta_i^L = \frac{\partial J}{\partial z_i^L} = \boxed{\frac{\partial J}{\partial y_i}} \frac{\partial y_i}{\partial z_i^L}$$


$\frac{\partial J}{\partial y_i}$  depends on the loss function

$$\delta^L = \nabla J(y) \odot \nabla a(z^L)$$

$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

(1) Initialization: compute  $\delta^L$


(2) Compute  $\delta^l$  based on  $\delta^{l+1}$



$$\begin{aligned} \delta_i^l &= \frac{\partial J}{\partial z_i^l} = \sum_k \left( \frac{\partial J}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial a_i^l} \frac{\partial a_i^l}{\partial z_i^l} \right) \\ &= \frac{\partial a_i^l}{\partial z_i^l} \sum_k \left( \frac{\partial J}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial a_i^l} \right) \end{aligned}$$

$\delta_i^{l+1}$

A red arrow points from the yellow box containing  $\frac{\partial J}{\partial z_k^{l+1}}$  to  $\delta_i^{l+1}$ .




$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

(1) Initialization: compute  $\delta^L$

(2) Compute  $\delta^l$  based on  $\delta^{l+1}$



$$\begin{aligned}\delta_i^l &= \frac{\partial a_i^l}{\partial z_i^l} \sum_k \left( \frac{\partial J}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial a_i^l} \right) \\ &= \frac{\partial a_i^l}{\partial z_i^l} \sum_k \delta_k^{l+1} w_{ki}^{l+1} \\ &= \nabla a_i^l(z_i^l) \sum_k \delta_k^{l+1} w_{ki}^{l+1}\end{aligned}$$



$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

Rethink the propagation

$$\delta_i^l = \nabla a_i^l(z_i^l) \sum_k \delta_k^{l+1} w_{ki}^{l+1}$$



$$\delta^l = \nabla a(z^l) \odot (W^{l+1})^T \delta^{l+1}$$

$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$

$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$


Idea: from L to 1

(1) Initialization: compute  $\delta^L$

$$\delta^L = \nabla J(y) \odot \nabla a(z^L)$$

(2) Compute  $\delta^l$  based on  $\delta^{l+1}$

$$\delta^l = \nabla a(z^l) \odot (W^{l+1})^T \delta^{l+1}$$






# Backpropagation

$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} = \begin{cases} a_j^{l-1} & , l > 1 \\ x_j & , l = 1 \end{cases}$$

## Forward Pass

$$\begin{aligned} z^1 &= W^1 x + b^1 & a^1 &= \sigma(z^1) \\ \vdots & & & \\ z^l &= W^l a^{l-1} + b^l & a^l &= \sigma(z^l) \\ \vdots & & & \end{aligned}$$



# Backpropagation

$$\frac{\partial J(\theta)}{\partial w_{ij}^l} = \frac{\partial J(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial J(\theta)}{\partial z_i^l} = \delta_i^l$$


## Backward Pass

$$\delta^L = \nabla J(y) \odot \nabla a(z^L)$$

$$\delta^{L-1} = \nabla a(z^L) \odot (W^{L+1})^T \delta^{L+1}$$

$$\vdots$$

$$\delta^l = \nabla a(z^l) \odot (W^{l+1})^T \delta^{l+1}$$

$$\vdots$$


# Reading Materials

- Automatic Differentiation in Machine Learning: a Survey (2015)

# Summary: How to Train Multilayer Neural Nets?

- Define the loss function  $l(,)$  properly
- A procedure to compute loss  $l(,)$  (*forward propagation*)
- A procedure to compute gradient  $\nabla l(,)$  (*back propagation*)
- Regularizer and its gradient  $\Omega(,)$  and  $\nabla \Omega(,)$
- Perform gradient based optimization method

# Summary

- Model Architectures
  - Artificial neurons
  - Activation function and saturation
  - Feedforward neural nets
- How to train a neural net
  - Loss Function Design
  - Optimization
    - Gradient Descent and Stochastic Gradient Descent
    - Back-propagation